

# Let's Get Started

```
library(tidyverse)
library(timetk)

# Setup for the plotly charts (# FALSE returns ggplots)
interactive <- TRUE
```

## Correlation Plots

`plot_acf_diagnostics()` returns the ACF and PACF of a target and optionally CCF's of one or more lagged predictors in interactive plotly plots. We also scale to multiple time series using `group_by()`.

- ACF = Autocorrelation between a target variable and lagged versions of itself.
- PACF = Partial Autocorrelation removes the dependence of lags on other lags highlighting key seasonalities.
- CCF = Shows how lagged predictors can be used for prediction of a target variable.

### Lag Specification

Lags (`.lags`) can either be specified as:

- A time-based phrase indicating a duration (e.g. 2 months)
- A maximum lag (e.g. `.lags = 28`)
- A sequence of lags (e.g. `.lags = 7:28`)

### Scales to Multiple Time Series with Groups

The `plot_acf_diagnostics()` works with grouped\_df's, meaning you can group your time series by one or more categorical columns with `dplyr::group_by()` and then apply `plot_acf_diagnostics()` to return group-wise lag diagnostics.

### Special Note on Groups

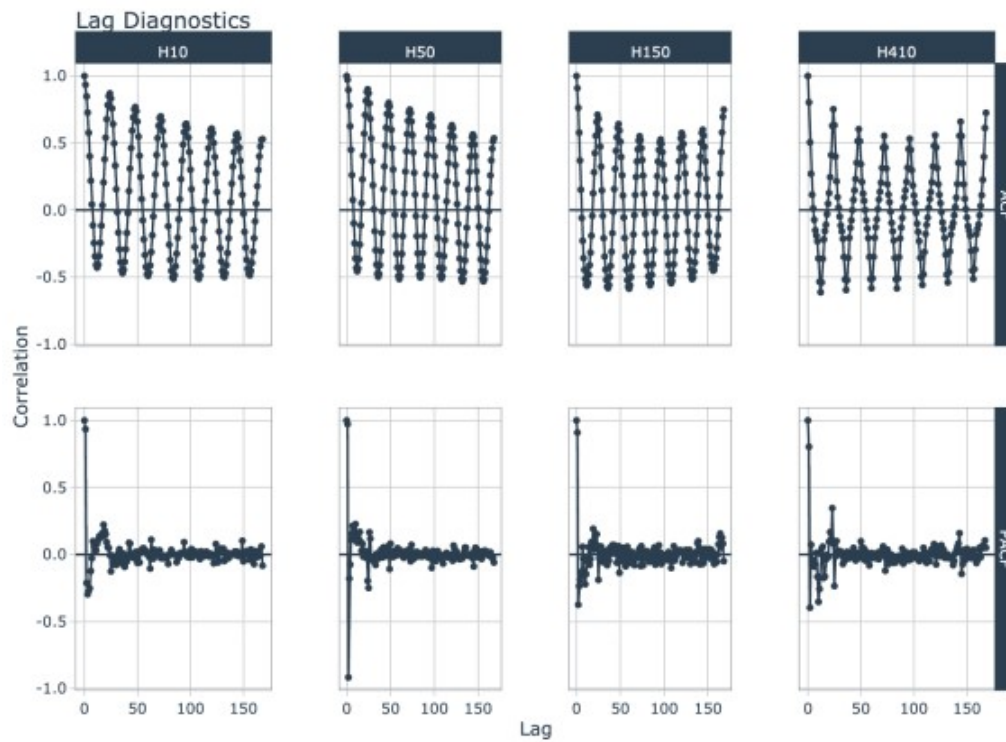
Unlike other plotting utilities, the `.facet_vars` arguments is NOT included. Use `dplyr::group_by()` for processing multiple time series groups.

### Calculating the White Noise Significance Bars

The formula for the significance bars is  $\pm 2/\sqrt{T}$  and  $-2/\sqrt{T}$  where  $T$  is the length of the time series. For a white noise time series, 95% of the data points should fall within this range. Those that don't may be significant autocorrelations.


## Grouped ACF Diagnostics

```
m4_hourly %>%
  group_by(id) %>%
  plot_acf_diagnostics(
    date, value,                # ACF & PACF
    .lags = "7 days",          # 7-Days of hourly lags
    .interactive = interactive
  )
```



## Grouped CCF Plots

```
walmart_sales_weekly %>%
  select(id, Date, Weekly_Sales, Temperature, Fuel_Price) %>%
  group_by(id) %>%
  plot_acf_diagnostics(
    Date, Weekly_Sales,          # ACF & PACF
    .ccf_vars = c(Temperature, Fuel_Price), # CCFs
    .lags      = "3 months",      # 3 months of weekly lags
    .interactive = interactive
  )
```



## Seasonality

`plot_seasonal_diagnostics()` is an interactive and scalable function for visualizing time series seasonality.

## Automatic Feature Selection

Internal calculations are performed to detect a sub-range of features to include using the following logic:

- The *minimum* feature is selected based on the median difference between consecutive timestamps
- The *maximum* feature is selected based on having 2 full periods.

Example: Hourly timestamp data that lasts more than 2 weeks will have the following features: “hour”, “wday.lbl”, and “week”.

## Scalable with Grouped Data Frames

This function respects grouped data.frame and tibbles that were made with `dplyr::group_by()`.

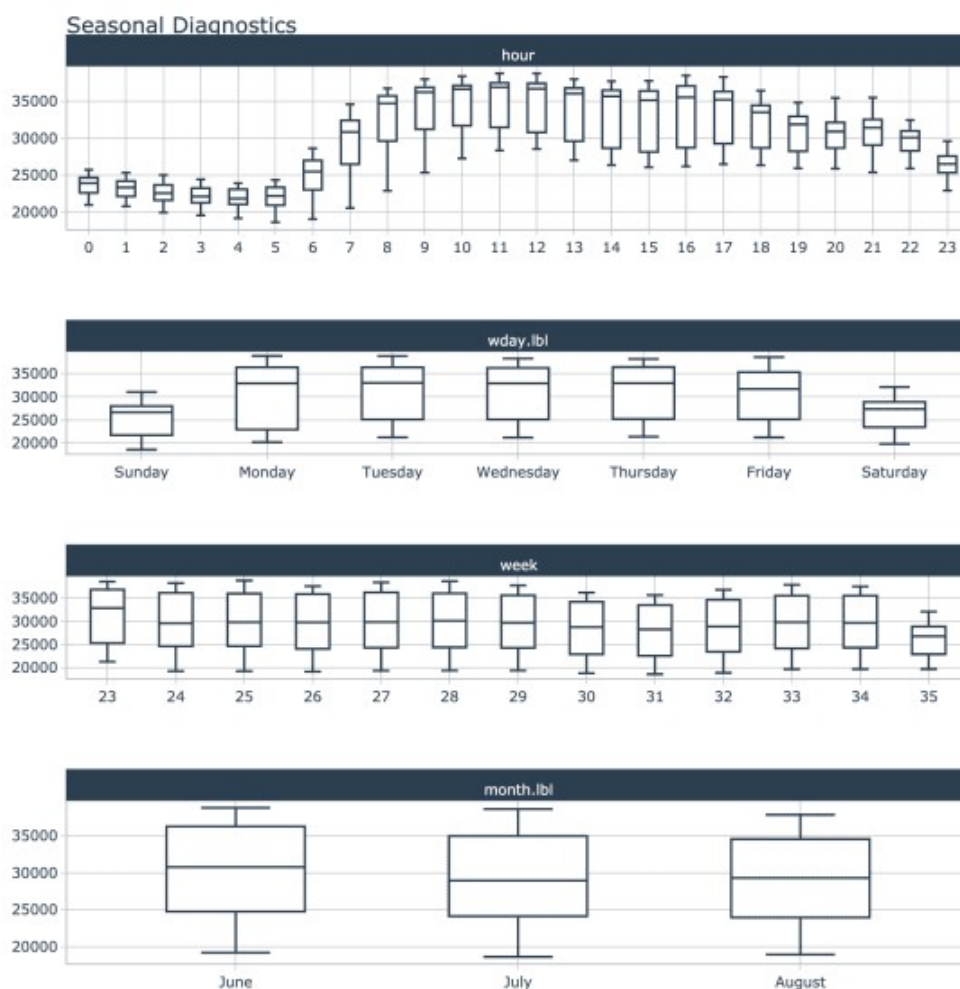
For grouped data, the automatic feature selection returned is a collection of all features within the sub-groups. This means extra features are returned even though they may be meaningless for some of the groups.

## Transformations

The `.value` parameter respects transformations (e.g. `.value = log(sales)`)

## Seasonal Visualizations

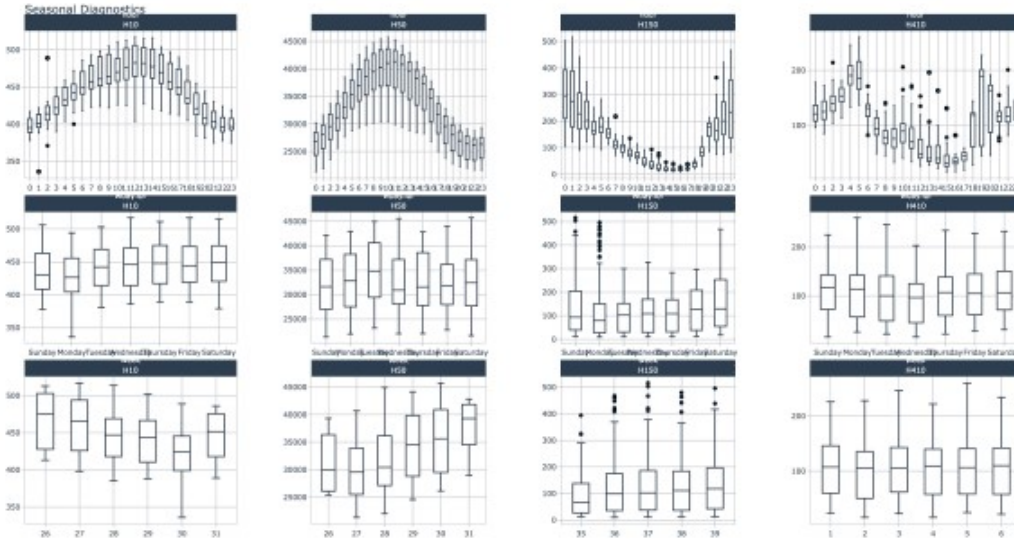
```
taylor_30_min %>%  
  plot_seasonal_diagnostics(date, value, .interactive = interactive)
```



## Grouped Seasonal Visualizations

```
m4_hourly %>%
```

```
group_by(id) %>%
plot_seasonal_diagnostics(date, value, .interactive = interactive)
```



## STL Diagnostics

The `plot_stl_diagnostics()` function generates a Seasonal-Trend-Loess decomposition. The function is “tidy” in the sense that it works on data frames and is designed to work with `dplyr` groups.

### STL method

The STL method implements time series decomposition using the underlying `stats::stl()`. The decomposition separates the “season” and “trend” components from the “observed” values leaving the “remainder”.

### Frequency & Trend Selection

The user can control two parameters: `.frequency` and `.trend`.

- The `.frequency` parameter adjusts the “season” component that is removed from the “observed” values.
- The `.trend` parameter adjusts the trend window (`t.window` parameter from `stl()`) that is used.

The user may supply both `.frequency` and `.trend` as time-based durations (e.g. “6 weeks”) or numeric values (e.g. 180) or “auto”, which automatically selects the frequency and/or trend based on the scale of the time series.

```
m4_hourly %>%
  group_by(id) %>%
  plot_stl_diagnostics(
    date, value,
    .frequency = "auto", .trend = "auto",
    .feature_set = c("observed", "season", "trend", "remainder"),
    .interactive = interactive)
```

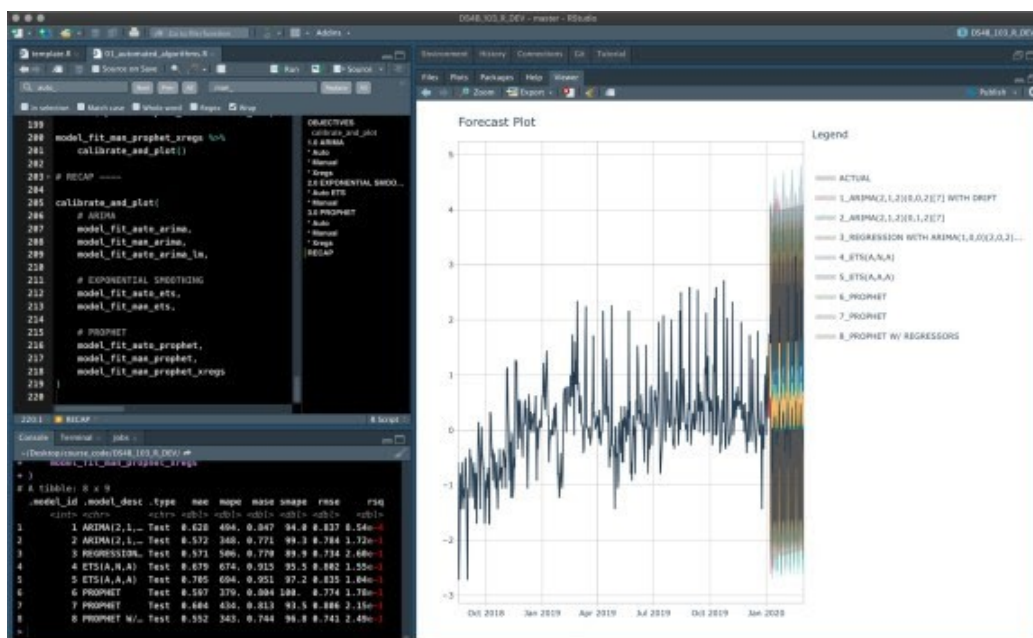


## Advanced Time Series Course

Become the times series domain expert in your organization.

Make sure you're notified when my new **Advanced Time Series Forecasting in R** course comes out. You'll learn `timetk` and `modeltime` plus the most powerful time series forecasting techniques available. Become the times series domain expert in your organization.

🔔 Get notified here: [Advanced Time Series Course](#).



You will learn:

- Time Series Preprocessing, Noise Reduction, & Anomaly Detection
- Feature engineering using lagged variables & external regressors
- Hyperparameter tuning
- Time series cross-validation
- Ensembling Multiple Machine Learning & Univariate Modeling Techniques (Competition Winner)
- NEW – Deep Learning with RNNs (Competition Winner)

- and more.